

# MongoDB

3IF Lab Support

Előd EGYED-ZSIGMOND

# Plan

Introduction

Data modeling

Practice

CRUD queries

Aggregates

Indexes

Replication

Sharding/ Partitioning

Conclusion

# Introduction

MongoDB, one of the most popular "NoSQL" DBMS

- It is a "document based" NoSQL system (with CouchDB, ElasticSearch, ...)
- relies on a semi-structured data model of (JSON encoding);
- Schema less (flexibility);
- An original (and specific) **query language**;
- No (or very little) of **transactional support**.

Built from the beginning as a **scalable** and **distributed** system

- Distribution by partitioning (sharding) ;
- Fault tolerance through replication.

# Why MongoDB

---

- A MongoDB server manages **data bases**.
- A database contains **collections**.
- The collections have **documents**.
- Each document has a unique identifier generated by MongoDB, **\_id** field.

# Document oriented

- The documents (lines) of the same kind are stored in collections (tables)
- A document is a tree, made of keys and values
- A value can be:
  - Scalar (int, long, string, date, binary, bool, etc.)
  - Array
  - A nested document

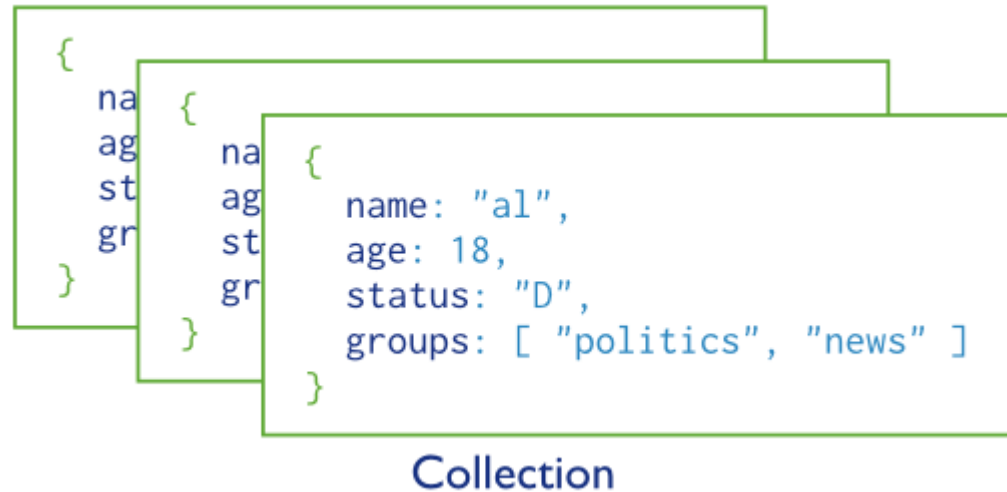
# Vocabulary

- Document

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

← field: value  
← field: value  
← field: value  
← field: value

- Collection



# Plan

Introduction

Data modeling

Practice

CRUD queries

Aggregates

Indexes

Replication

Sharding/ Partitioning

Conclusion

# How to model

- complex document
  - nested documents
  - the value objects are part of the document
  - Sometimes greedy (denormalisation)
  - (current) limit: 4MB per document
- normalization
- mixed approach



# JSON: Modeling

## Relational


Person:

Pers_ID	Surname	First_Name	City
0	Miller	Paul	London
1	Ortega	Alvaro	Valencia
2	Huber	Urs	Zurich
3	Blanc	Gaston	Paris
4	Bertolini	Fabrizio	Rom

Car:

Car_ID	Model	Year	Value	Pers_ID
101	Bentley	1973	100000	0
102	Rolls Royce	1965	330000	0
103	Peugeot	1993	500	3
104	Ferrari	2005	150000	4
105	Renault	1998	2000	3
106	Renault	2001	7000	3
107	Smart	1999	2000	2

no relation

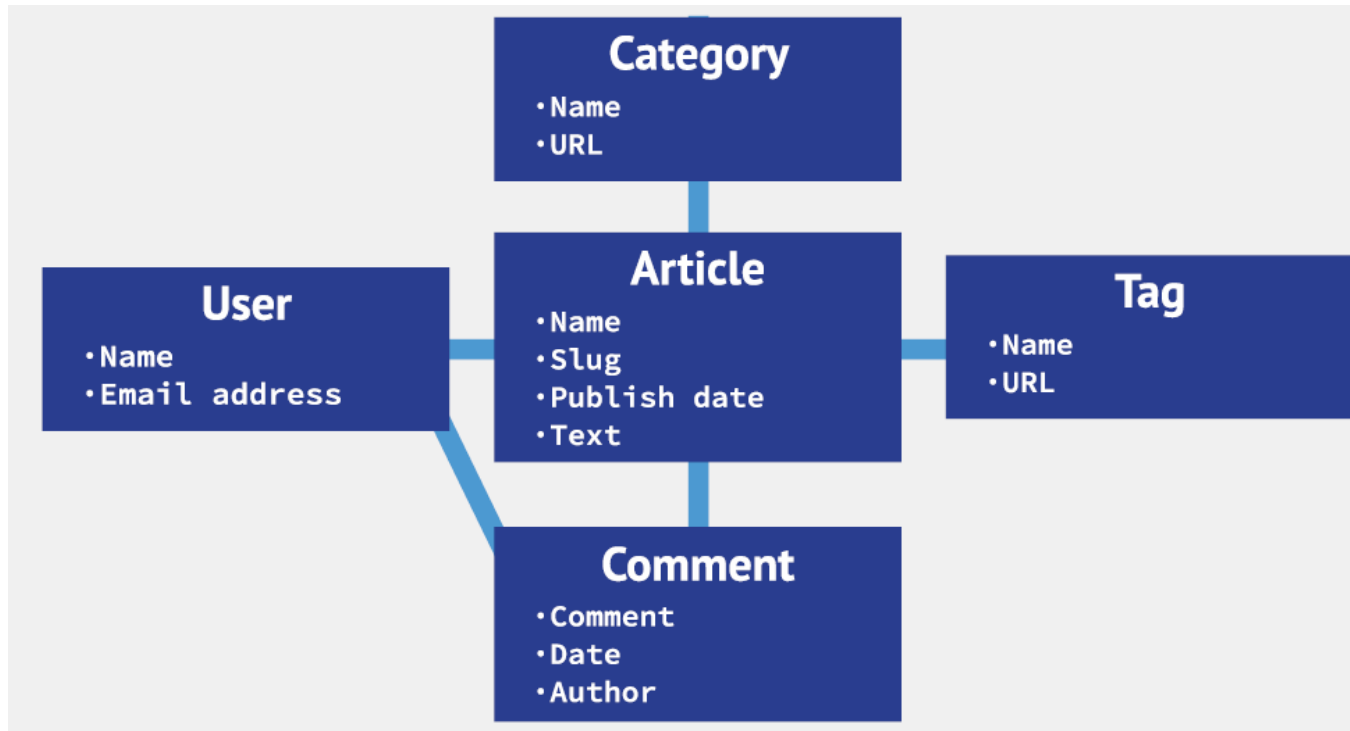


## MongoDB Document

```
{
  first_name: 'Paul',
  surname: 'Miller'
  city: 'London',
  location: [45.123,47.232],
  cars: [
    { model: 'Bentley',
      year: 1973,
      value: 100000, ... },
    { model: 'Rolls Royce',
      year: 1965,
      value: 330000, ... }
  ]
}
```

Mat Keep: <https://www.slideshare.net/matkeep/migrating-from-relational-databases-to-mongodb>

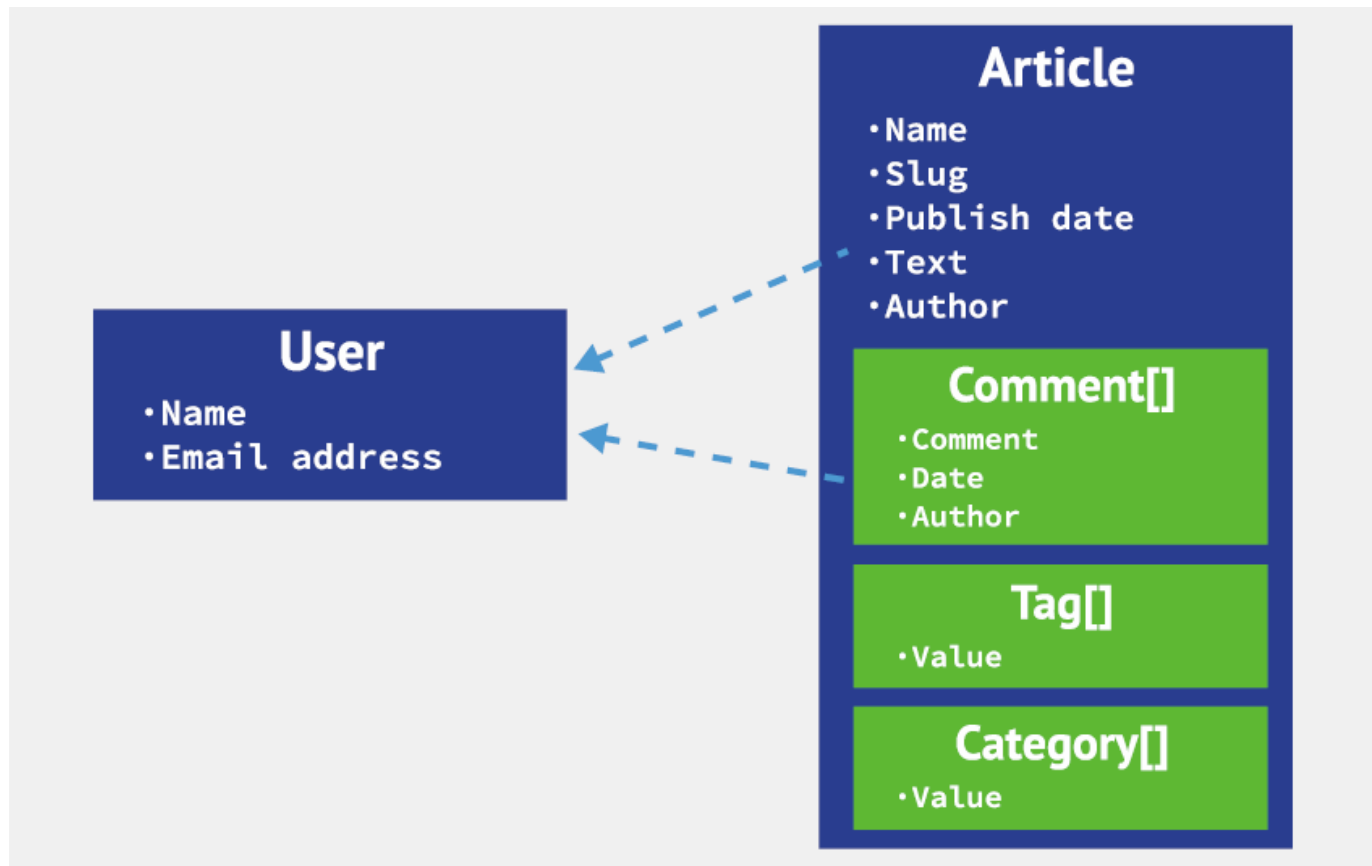
# JSON: Modeling



RDBMS: Join 5 tables

Mat Keep: <https://www.slideshare.net/matkeep/migrating-from-relational-databases-to-mongodb>

# JSON: Modeling



MongoDB: Denormalized to 2 BSON documents

Mat Keep: <https://www.slideshare.net/matkeep/migrating-from-relational-databases-to-mongodb>

# Document Model Benefits

- Rich data model, natural data representation
  - Embed related data in sub-documents & arrays
  - Support indexes and rich queries against any element
- Data aggregated to a single structure (pre-JOINED)
  - Programming becomes simple
  - Performance can be delivered at scale
- Dynamic schema
  - Data models can evolve easily
  - Adapt to changes quickly: agile methodology

# Plan

Introduction

Data modeling

Practice

CRUD queries

Aggregates

Indexes

Replication

Sharding/ Partitioning

Conclusion

# Convenient

MongoDB comes with a shell: `bin/mongo`

Start server with: `bin/mongod`

Some arguments:

- `--dbpath <path>`: Data Storage path
- `--port <port>`: Port the server listens
- `--replSet <Name>`: Enter the server in a replicas cluster

# Convenient

MongoDB comes with a shell: `bin/mongosh`

Start server with: `bin/mongod`

Some arguments:

- `--dbpath <path>`: Data Storage path
- `--port <port>`: Port the server listens
- `--replSet <Name>`: Enter the server in a replicas cluster

# Shell

- View the current database:
  - `db`
- See list of existing databases:
  - `show dbs`
- Select / create a database:
  - `use <name>`
- View Collections:
  - `show collections`
- Create collection (required)
  - `db.createCollection`



# Plan

Introduction

Data modeling

Practice

CRUD queries

Aggregates

Indexes

Replication

Sharding/ Partitioning

Conclusion

# Queries in MongoDB

---

- JSON queries
- We find the equivalent of SQL projections
- Some default aggregations
  - ... Map/Reduce for the others

# SQL to Aggregation Mapping

SQL Terms, Functions, and Concepts	MongoDB Aggregation Operators
WHERE	\$match
GROUP BY	\$group
HAVING	\$match
SELECT	\$project
ORDER BY	\$sort
LIMIT	\$limit
SUM()	\$sum
COUNT()	\$sum \$sortByCount
join	\$lookup

Mapping Chart:

<http://docs.mongodb.org/manual/reference/sql-aggregation-comparison/>

# CRUD

## Modification:

- Insert
- Update
- Remove

The operations are applied to a collection

<http://docs.mongodb.org/manual/>

# CRUD

## Insertion

Collection  
↓  
`db.users.insert(`

Document  
↓

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

)

<http://docs.mongodb.org/manual/>

Document

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

insert

Collection

{ name: "al", age: 18, ... }
{ name: "lee", age: 28, ... }
{ name: "jan", age: 21, ... }
{ name: "kai", age: 38, ... }
{ name: "sam", age: 18, ... }
{ name: "mel", age: 38, ... }
{ name: "ryan", age: 31, ... }
{ name: "sue", age: 26, ... }

users

The "\_id" field is automatically added / alternative to save

# CRUD

## Update

```
db.collection.update(  
  <Query>           Same constraints as find  
  <Update>          $set, $unset...  
  {  
    upsert<boolean>  
    Multi <boolean>  
    writeConcern: <Document>  
  }  
)
```

```
db.users.update(  
  { age: { $gt: 18 } },  
  { $set: { status: "A" } },  
  { multi: true }  
)
```

← collection  
← update criteria  
← update action  
← update option

Default update concerns a single document

<http://docs.mongodb.org/manual/>

# CRUD

## Update

```
db.movies.update({ "year":{$lt: 2000}}, {$set: {old: true}})
```

```
db.movies.find({old: true})
```

```
db.movies.updateMany({ "year":{$lt: 2000}}, {$set: {old: true}})
```

*(since  
MongoDB 3.2)*

```
db.movies.update({ "year":{$lt: 2000}}, {$unset : {old: true}})
```

```
db.movies.update({ "year":{$lt: 2000}}, {$set: {old: true}},  
{multi: true})
```

# CRUD

## Update, update a table

- We will add a table to a document

```
db.produits.insert({Counter: 100001, tab: [ 'a', 'b', 'c']})
```

- You can view it with the following command:

```
db.produits.find({Counter: 100001})
```

- Now to add an item to the table, use the operand **\$push**:

```
db.produits.update({Counter: 100001}, {$push: {tab: 'd'}})
```

To add several items at once, there is **\$pushAll** to add a whole table.



# CRUD

## Update a table

- The **\$pop** operand will remove the last item:

```
db.produits.update({Counter: 100001} {$pop: {tab: 1}})
```

- The array has lost its last element "d".
- To remove the first element:

```
db.produits.update({Counter: 100001} {$pop: {tab: -1}})
```

- Analogously with the sort method, with -1 elements are deleted in the other direction.
- With the operand **\$addToSet** an item is added without duplication

```
db.produits.update({Counter: 100001} {$addToSet : {Tab: 'b'}})
```

- The "b" value was present in the table so it is not added.

# CRUD

## Changes, isolation

```
db.films.update(  
  { old: true, $isolated : 1},  
  {$inc: {count: 1}},  
  {multi: true }  
)
```

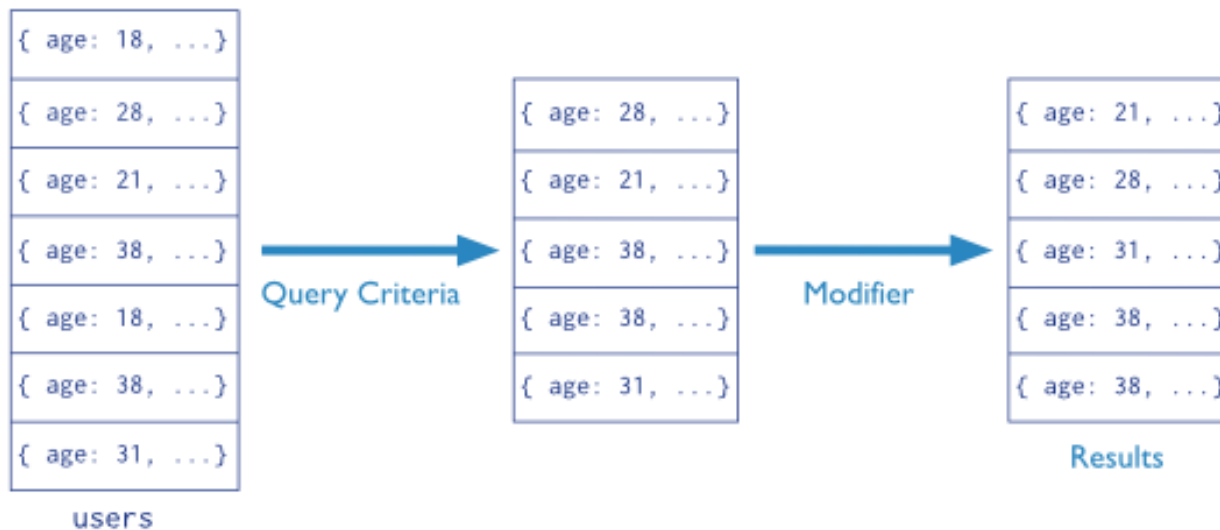
```
db.products.remove(  
  { qty: {$gt: 20}, $isolated: 1})
```

\$isolated does not work on distributed collections

# CRUD

## queries

Collection                      Query Criteria                      Modifier  
`db.users.find( { age: { $gt: 18 } } ).sort( {age: 1 } )`



<http://docs.mongodb.org/manual/>

# CRUD

## queries

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
) .limit(5)
```

← collection  
← query criteria  
← projection  
← cursor modifier

## in SQL

```
SELECT _id, name, address  
FROM   users  
WHERE  age > 18  
LIMIT 5
```

← projection  
← table  
← select criteria  
← cursor modifier

```
db.movies.find({ "country": {$ne: "USA"}}, {_id: 0, title: 1  
}).limit(5)
```

<http://docs.mongodb.org/manual/>

# operators

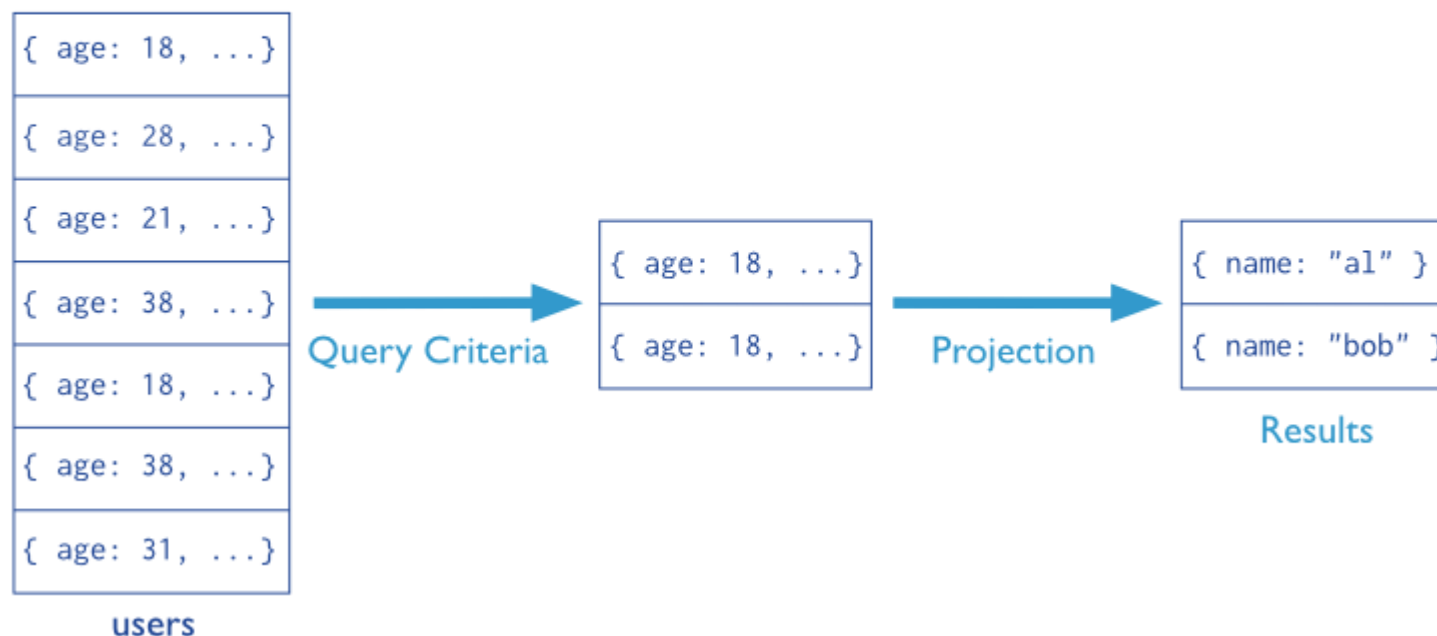
\$eq	=	<b>values which are equal to a specified value.</b>
\$gt	>	values that are greater than a specified value.
\$gte	> =	values that are greater than or equal to a specified value.
\$lt	<	values which are less than a specified value.
\$lte	<=	values that are less than or equal to a specified value.
\$ne	!=	all values that are not equal to a specified value.
\$in	∈	one of the values specified in a table.
\$nin	∉	none of the values specified in the table.
\$or		
\$and		
\$not		
\$nor		
...		

<http://docs.mongodb.org/manual/reference/operator/query/>

# CRUD

## projection

**Collection**                      **Query Criteria**                      **Projection**  
`db.users.find( { age: 18 }, { name: 1, _id: 0 } )`



```
db.movies.find({ "country":{$ne: "USA"}}, {_id: 0, title: 1})
```

# CRUD Examples

Exclude the year:

```
db.films.find({ "country":"USA"}, {year: 0})
```

Return the title and genre and `_id`

```
db.films.find({ "country":"USA"}, {title: 1, type: 1})
```

Return movies with a given role

```
db.films.find({ "actors.role":"William Munny"})
```

Return the title and genre

```
db.films.find({ "country":" USA"}, {title: 1, type: 1, _id: 0})
```

Sort Results

```
db.films.find({ "year":{$gt: 2000}}).sort({year: -1})
```

# CRUD

## Examples

```
db.movies.find({$and : [{"year": 2003}, {genre: "romance"}] },{"title":1,"genre":1});
```

```
db.movies.find({$or : [{"year": 2003}, {genre: "romance"}] },{"title":1,"genre":1});
```



# CRUD

## sliders

```
db.collection.find()
```

```
db.films.find({ "country": "FR" })
```

## Course one by one results

```
var myCursor = db.films.find( { country: 'FR' } );  
  
while (myCursor.hasNext()) {  
    print(tojson(myCursor.next()));  
}
```

```
var myCursor = db.films.find( { country: 'FR' } );  
|  
myCursor.forEach(printjson);
```

# CRUD

## Cursors, methods

- *limit(n)*: To retrieve only first n results
- *sort (...)*: To sort the results
- *skip (n)*: To skip n results

For example, for the penultimate film in the name, you could do:

```
var cur = db.films.find();  
cur.sort({title: -1}).limit(1).skip (2);
```

# CRUD

## Cursors

### Transformation into an array

```
var myCursor = db.films.find( { country: 'FR' } );  
var documentArray = myCursor.toArray();  
var myDocument = documentArray[1];  
  
print (toJson(myDocument));
```

## Analyze query performance

```
db.films.find({Country: 'FR'}).explain();  
db.films.createIndex({country: 1})
```

# Plan

Introduction

Data modeling

Practice

CRUD queries

Aggregates

Indexes

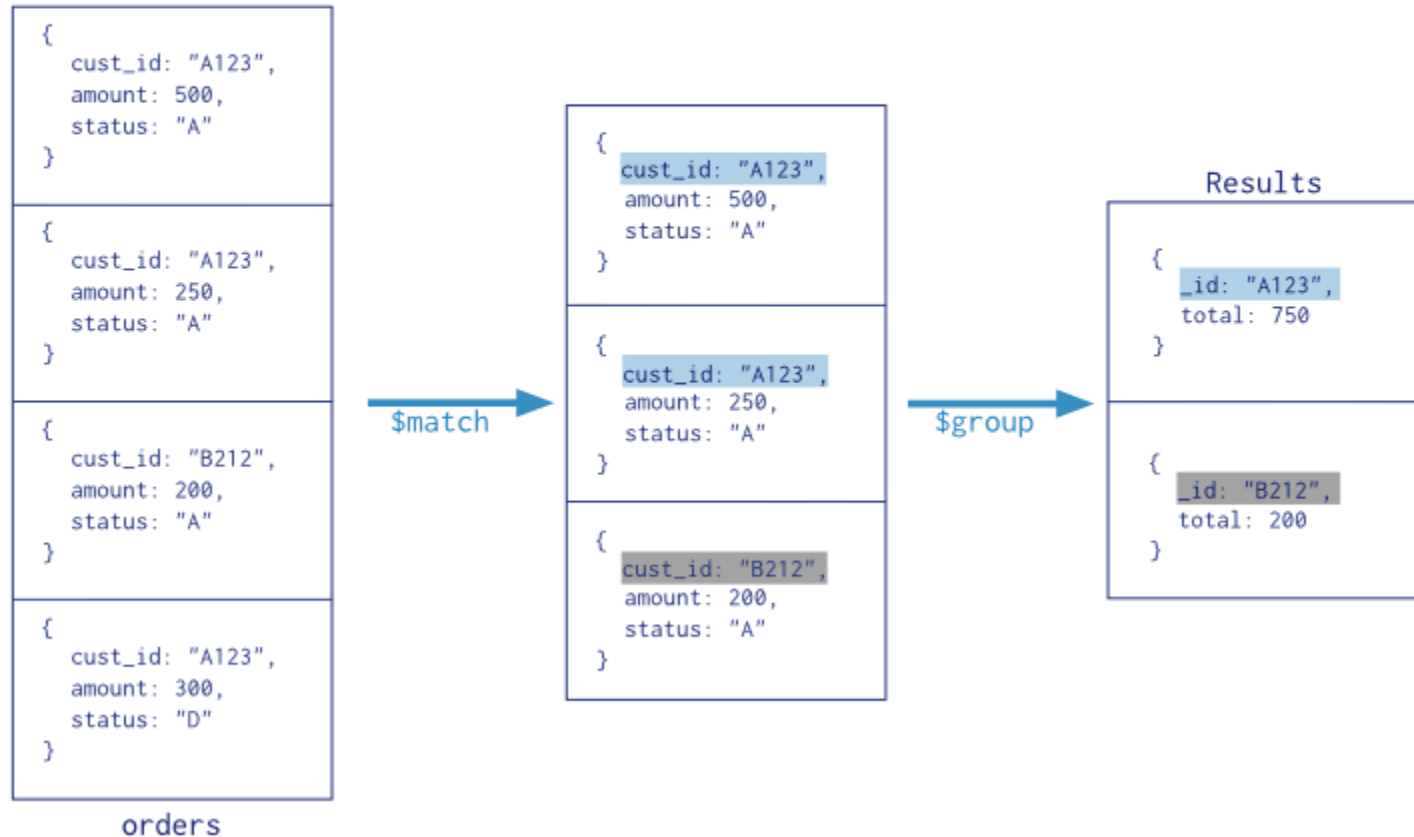
Replication

Sharding/ Partitioning

Conclusion

# aggregates

Collection  
↓  
`db.orders.aggregate( [`  
    `$match stage → { $match: { status: "A" } },`  
    `$group stage → { $group: { _id: "$cust_id", total: { $sum: "$amount" } } }`  
    `]` `)`



# aggregates

## Nb films by genre

```
db.movies.aggregate([
  {$group: {_id: "$genre", count: {$sum: 1}}}
])
```

## Number of American films by genre

```
db.movies.aggregate([
  {$match: {country: "USA"}},
  {$group: {_id: "$genre", count: {$sum: 1}}}
])
```

## Date of first horror movie

```
db.movies.aggregate([
  {$match: {genre: "Horreur"}},
  {$group: {_id: "$genre", debut: {$min: "$year"}}}
])
```

# Plan

Introduction

Data modeling

Practice

CRUD queries

Aggregates

Indexes

Replication

Sharding/ Partitioning


Conclusion

# indexes

- Very similar to the RDBMS, indexation in MongoDB is made on one or more fields.
- Improves search performance.
- The indexes are stored in the collections.
- Provides overload for write operations.
- The internal operation is very similar to that found in the current DBMS.



# indexes



<https://docs.mongodb.org>

# indexes

## Types of indexes

- single field
- composed field
- multi key
- geo spatial
- text
- Hashed

# indexes

- Consider using indexes effectively.
- A query, a bit field has no interest in being indexed
- Although we speak of NoSQL the indexes are similar to those in the RDBMS.

# Loading Documents

- unitary insert:

- Example: `coll.insert({x:1,y:2,z:"test"})`
- Command to avoid for large inserts via a client (> 10 000 documents)

- Bulk insert:

-Example:

```
var Documents = [];  
    for(var i = 1; i < 10 000; i++) {  
documents.push({x:i, y:i, z:"test"});  
        db.coll.insert(documents);  
    }
```

# Inserting massive documents

- Program to load a BSON file containing hundreds of GB or more in a MongoDB base: *mongorestore*
- Mongorestore: Method far more effective to insert data (other methods: ( bulk ) Insert, mongoimport )
- Backup / Restore tool of a MongoDB base ( Sharded ) with *mongodump*

# other features

- geographic queries
- Javascript running on the server
  - db.eval (~ Stored procedures)
  - map/Reduce (Aggregations)
  - db.system.js (functions)

# Plan

Introduction

Data modeling

Practice

CRUD queries

Aggregates

Indexes

Replication

Sharding/ Partitioning

Conclusion

# Replication

What brings replication?

- Redundancy
- Simplify tasks (backups, ...)
- Increase reading ability
- A replica set is a cluster of MongoDB instances.
- master / slave strategy
- There must ALWAYS be a single master.



# Replication

---

- The master to the slave replication is asynchronous.
- Synchronous: Blocking / Expensive / Strong consistency
- Asynchronous: No blocking / Data refreshing mandatory.

# Plan

Introduction

Data modeling

Practice

CRUD queries

Aggregates

Indexes

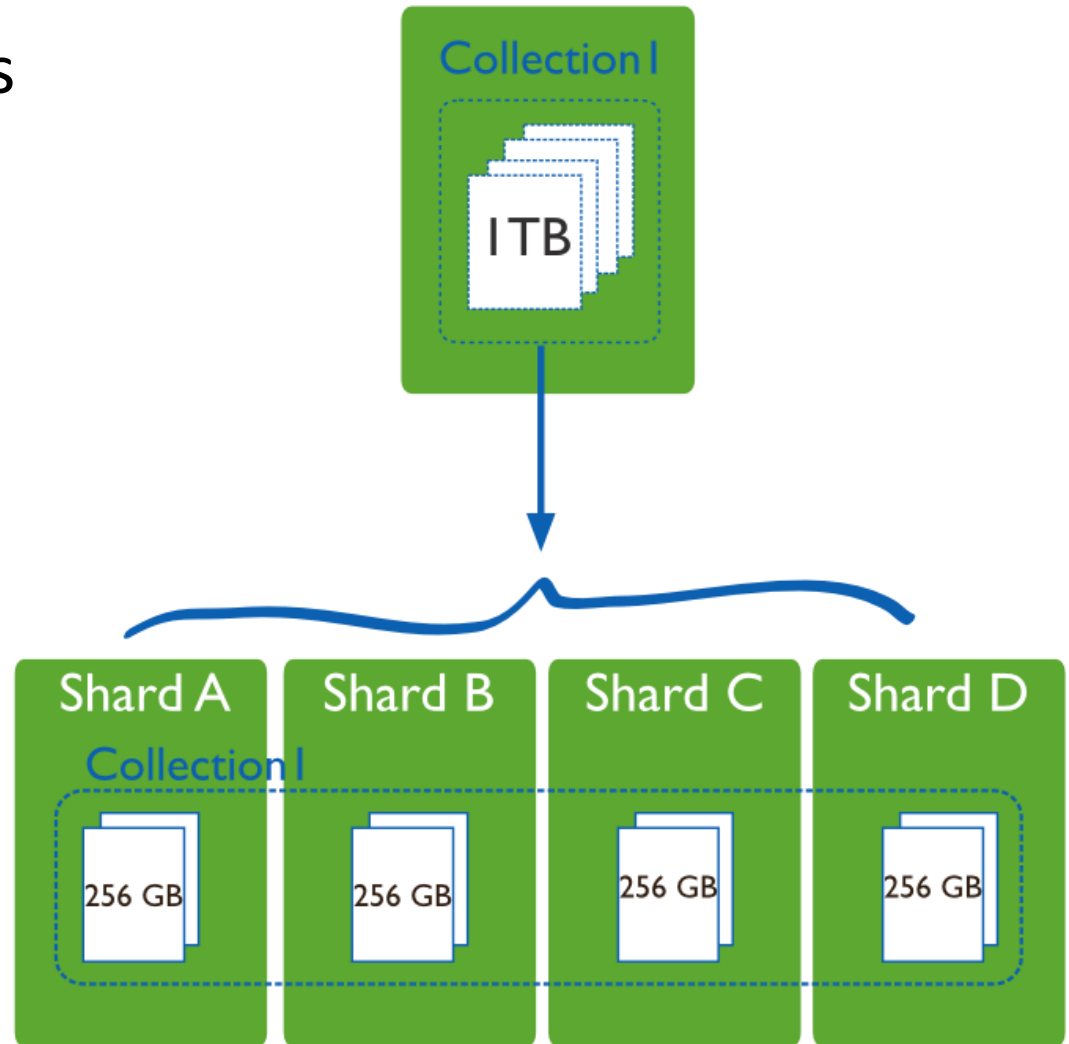
Replication

Sharding/ Partitioning

Conclusion

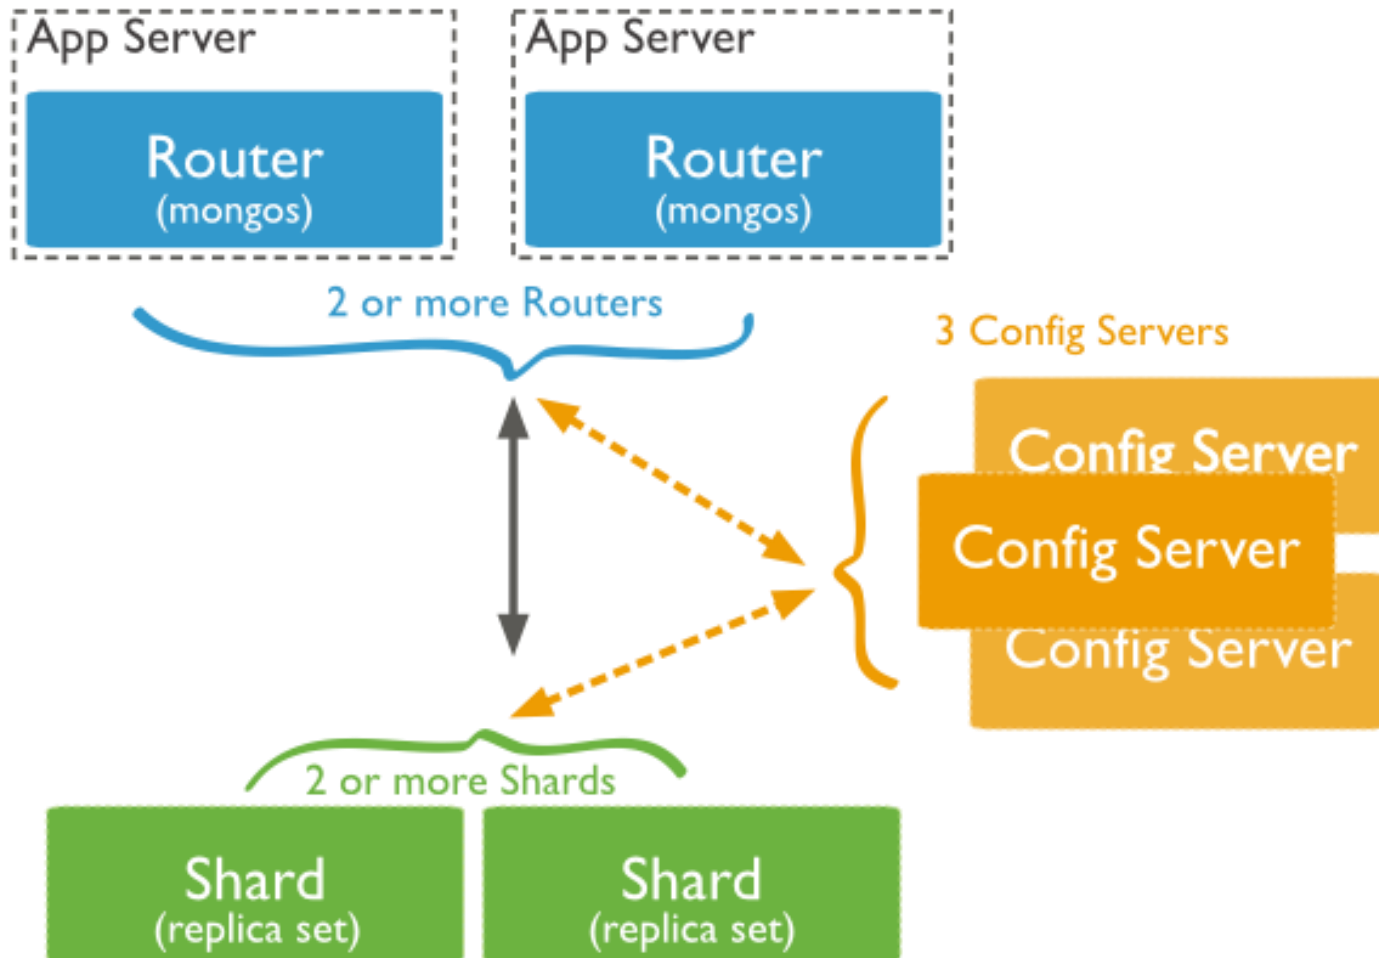
# sharding

- The sharding enables automatic data partitioning



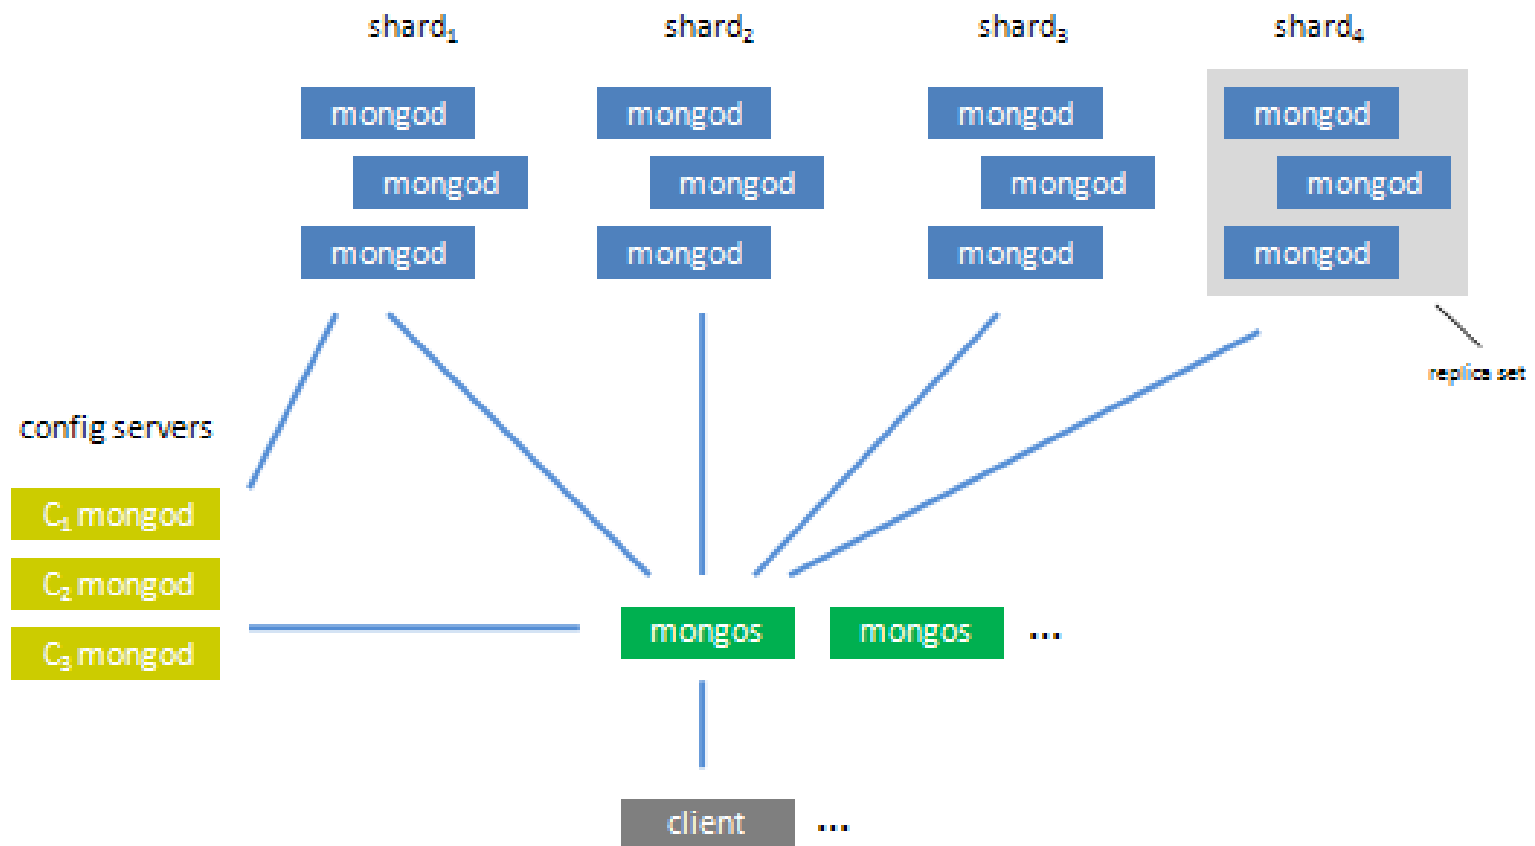
<https://docs.mongodb.org>

# sharding



<https://docs.mongodb.org>

# sharding and replication



# Plan

Introduction

Data modeling

Practice

CRUD queries

Aggregates

Indexes

Replication

Sharding/ Partitioning

Conclusion

# More Information

• Resource	Location
• MongoDB Downloads	<a href="https://mongodb.com/download">mongodb.com/download</a>
• Free Online Training	<a href="https://education.mongodb.com">education.mongodb.com</a>
• Webinars and Events	<a href="https://mongodb.com/events">mongodb.com/events</a>
• White Papers	<a href="https://mongodb.com/white-papers">mongodb.com/white-papers</a>
• Case Studies	<a href="https://mongodb.com/customers">mongodb.com/customers</a>
• Presentations	<a href="https://mongodb.com/presentations">mongodb.com/presentations</a>
• Documentation	<a href="https://docs.mongodb.org">docs.mongodb.org</a>

# References

<http://nosql.developpez.com/>

<http://news.humancoders.com/t/nosql/>

<https://docs.mongodb.com/manual/reference/sql-comparison/>

## MongoDB

Official site : [www.mongodb.org](http://www.mongodb.org)

A client: <https://studio3t.com/>